Emotion Discovery and Reasoning its Flip in Conversation(EDiReF)

Submitted in partial fulfillment of the requirements for the
award of the degree of

BACHELOR OF TECHNOLOGY
in
Computer Engineering

Submitted by:
Tirth B Dalwadi (20BCS051)
Mohammad Zohair (20BCS038)

Faculty Supervisor:
Dr. Tanvir Ahmad

Department of Computer Engineering
Faculty of Engineering and Technology
Jamia Millia Islamia – New Delhi- 110025

CERTIFICATE

This is to certify that the content of this project entitled "**Emotion Discovery and Reasoning its Flip in Conversation (EDiReF)**" by **Mohammad Zohair (20BCS038)**, and **Tirth B Dalwadi (20BCS051)** is a record of bonafide work carried out by them, in the **Department of Computer Engineering, Jamia Millia Islamia, New Delhi**, under my supervision and guidance in partial fulfillment of requirements for the award of Bachelor Of Technology in Computer Engineering, Jamia Millia Islamia in the academic year 2024.

Prof Bashir Alam
(Head of Department)
Department of Computer Engineering
Faculty of Engineering and Technology
Jamia Millia Islamia
New Delhi

Prof Tanvir Ahmad
(Project Supervisor)
Department of Computer Engineering
Faculty of Engineering and Technology
Jamia Millia Islamia
New Delhi

ACKNOWLEDGEMENT

We extend our heartfelt appreciation to Dr. Tanvir Ahmad, Professor in the Department of Computer Engineering at Jamia Millia Islamia, New Delhi, for his invaluable technical guidance, innovative insights, and unwavering support throughout this project. Our sincere gratitude also goes to Prof. Bashir Alam, our esteemed Head of Department, for his invaluable encouragement and continuous support. We express our deepest gratitude to the Department of Computer Engineering and the entire faculty members for their exceptional teaching, guidance, and encouragement that significantly contributed to the success of this project. Furthermore, we would like to acknowledge and thank our classmates and friends for their valuable suggestions and consistent support whenever needed. We apologize for any inadvertent omissions in this acknowledgment.

Mohammad Zohair
(20BCS038)

Tirth B Dalwadi
(20BCS051)

Department of Computer Engineering
Faculty of Engineering & Technology
JAMIA MILLIA ISLAMIA
NEW DELHI

ABSTRACT

Emotions have always been an inseparable part of human communication, ranging from face-to-face conversations, speeches involving facial expressions, and even text-based conversations, which turn out to be one of the most prevalent mediums of passing information in the current scenario. They are always one of the major driving forces of the course of conversation, its outcomes, and consequences on the involved parties as well as the associated environment. This prime importance of emotions brings with them a crucial challenge of recognizing emotions accurately in a sequence of dialogues. The subtleties of emotion variation are dynamic in nature, due to the possibilities of affirmation, negation, sarcasm, and other related communication components existing in the set of dialogues. Moreover, the emotional tendencies of the involved speakers play an important role in determining the conversation results. Thus, in accordance with the same, the domain of emotion recognition and figuring out the triggers of emotion flip have found emergence in recent times and gained attention from researchers for developing mechanisms to accomplish the related objectives. This research work is oriented towards the tasks of emotion recognition in conversations (ERC) and emotion flip reasoning (EFR)*. In this research study, we propose a meticulous mechanism for recognizing emotions in a series of conversations. The system has been implemented to understand the speaker-oriented information, dynamics, and important features affecting the emotion-based situation of the conversation. In addition to this, we aim to propose a mechanism for detecting and extracting the triggers existing in the course of dialogues which lead to a significant change of emotional state of the involved parties. The proposed system architecture has been developed using pre-trained language models, including Multilingual BERT and DistilBERT to understand the context as well as the sentiments of the dialogues, in accordance with a code-mixed dataset consisting of conversations from a TV show.

Keywords: Emotion Recognition in conversations (ERC), Emotion Flip Recognition (EFR), Large Language Models, Sentiment Analysis, context-cues, speaker-aware mechanism

**This research work will be a part of the Shared Task-10 at the 18th International Workshop on Semantic Evaluation (SemEval 2024).*

TABLE OF CONTENTS

S. No.	Contents		Page No.
	CERTIFICATE		2
	ACKNOWLEDGEMENT		3
	ABSTRACT		4
	TABLE OF CONTENTS		5
1	INTRODUCTION		7
	1.1	Emotions	7
	1.2	Emotion Recognition	8
	1.3	Emotion Recognition in Conversation	9
	1.4	Challenges in Emotion Recognition in Conversation	10
	1.5	Translation	11
	1.6	Transliteration	12
2	LITERATURE REVIEW		13
3	RESEARCH GAPS		16
4	RESEARCH OBJECTIVES		16
5	THEORETICAL BACKGROUND		17
	5.1	Transformers	17
	5.2	Transfer Learning	17
	5.3	Architecture of Transformers	18
	5.4	Attention Layers	19
	5.5	Encoder Models	19

	5.6	Decoder Models	20
	5.7	Pre-trained Language Models	20
6	PROPOSED METHODOLOGY		22
	6.1	Dataset Used	22
	6.2	Preprocessed Dataset	23
	6.3	Implementation Strategy	24
	6.4	Model Architecture	26
	6.5	Model Training	27
7	EVALUATION METRICS		29
	7.1	DistilBERT BASE (UNCASED)	29
	7.2	BERT MULTILINGUAL BASE (CASED)	30
8	EVALUATION METRICS		31
	8.1	DistilBERT BASE (UNCASED)	31
	8.2	BERT MULTILINGUAL BASE (CASED)	32
9	FUTURE WORK		34
10	CONCLUSION		35
11	PROGRAMMING ENVIRONMENT AND TOOLS USED		35
12	REFERENCES		36

1. INTRODUCTION

1.1 Emotions

Fig.1

In "Discovering Psychology," authors Don Hockenbury and Sandra E. Hockenbury suggest that emotions are complex psychological states that involve three distinct components: a subjective experience, a physiological response, and a behavioral or expressive response.

Emotion is a complex psychological and physiological experience that involves a range of feelings, thoughts, and physiological responses. It is a subjective state of consciousness that is often triggered by specific events, thoughts, or external stimuli.

Emotions can influence behavior, decision-making, and interpersonal relationships. There are basic or primary emotions that are considered universal across cultures, such as happiness, sadness, anger, fear, surprise, and disgust. These emotions are thought to be innate and serve evolutionary purposes.

In the context of conversation, emotions play a crucial role in shaping the dynamics and outcomes of communication. When individuals engage in a conversation, they bring with them a spectrum of emotions that influence their thoughts, expressions, and responses. These emotions are often triggered by the content of the conversation, the tone of voice, body language, or specific cues.

During a conversation, people may experience a variety of emotions such as happiness, sadness, frustration, excitement, or empathy, depending on the topics discussed and the way the interaction unfolds. These emotional experiences can impact not only the verbal aspects of communication but also nonverbal cues, including facial expressions, gestures, and posture.

Emotions are typically accompanied by physiological changes, including changes in heart rate, facial expressions, and hormonal levels. The brain plays a crucial role in generating and regulating emotions, with different areas of the brain being associated with specific emotions.

1.2 Emotion Recognition

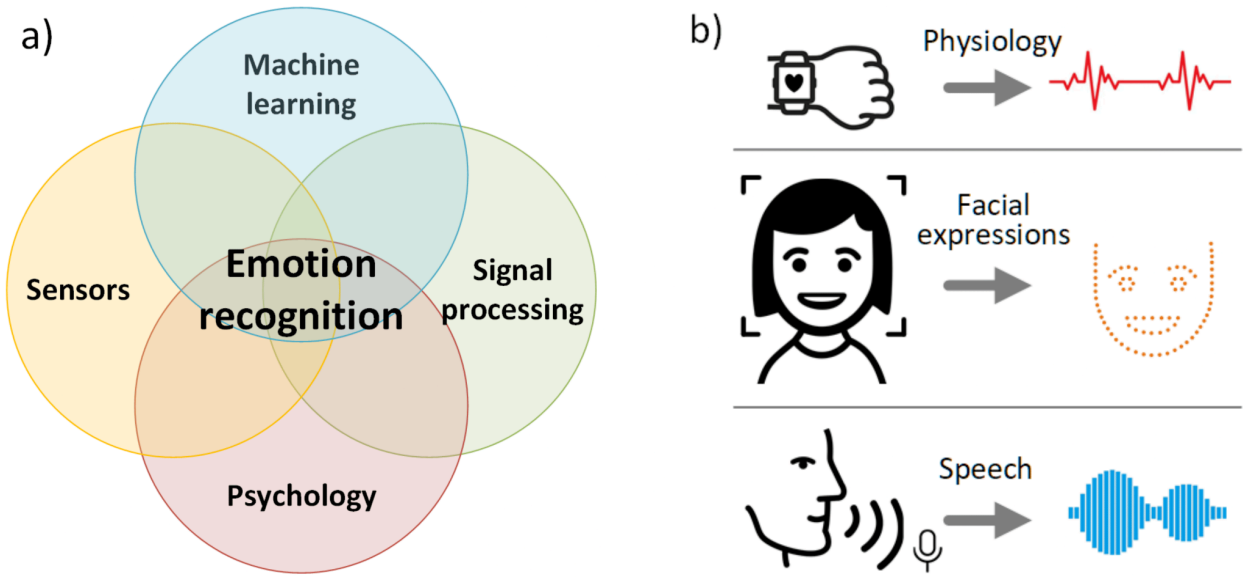

Fig. 2

Emotion recognition refers to the ability to identify, interpret, and understand the emotional states of individuals based on various cues, such as facial expressions, vocal tone, body language, and physiological signals. This process involves both perceptual and cognitive aspects, as individuals must not only detect emotional cues but also comprehend and attribute specific emotions to them.

In the realm of technology and artificial intelligence, emotion recognition often involves the development and implementation of algorithms and systems capable of analyzing human expressions and behaviors to infer emotional states. Some common modalities used in emotion recognition include:

- Facial Expression Analysis: This involves detecting and interpreting facial expressions to identify emotions such as happiness, sadness, anger, surprise, and more. Computer vision techniques and facial recognition algorithms are often employed for this purpose.
- Voice Analysis: Emotion recognition systems may analyze the tone, pitch, and other vocal characteristics to determine emotional states in speech. Speech processing technologies play a key role in identifying emotions conveyed through spoken words.
- Body Language and Gestures: Observation of body movements, gestures, and posture can provide valuable cues for recognizing emotions. This can be achieved through computer vision technologies that track and analyze body movements.
- Physiological Signals: Monitoring physiological indicators like heart rate, skin conductance, or EEG (electroencephalogram) can offer insights into emotional states.

These measures are often used in research and certain applications to complement other modalities.

Emotion recognition has applications in various fields, including human-computer interaction, customer service, mental health monitoring, and market research. In human-computer interaction, for example, systems may adapt their responses based on the user's emotional state to create more personalized and responsive experiences.

It's important to note that while significant progress has been made in the development of emotion recognition technologies, challenges remain, such as the cultural variability of expressions and the need for context-aware interpretation of emotions. Additionally, ethical considerations related to privacy and consent are important in the deployment of such technologies.

1.3 Emotion Recognition in Conversation

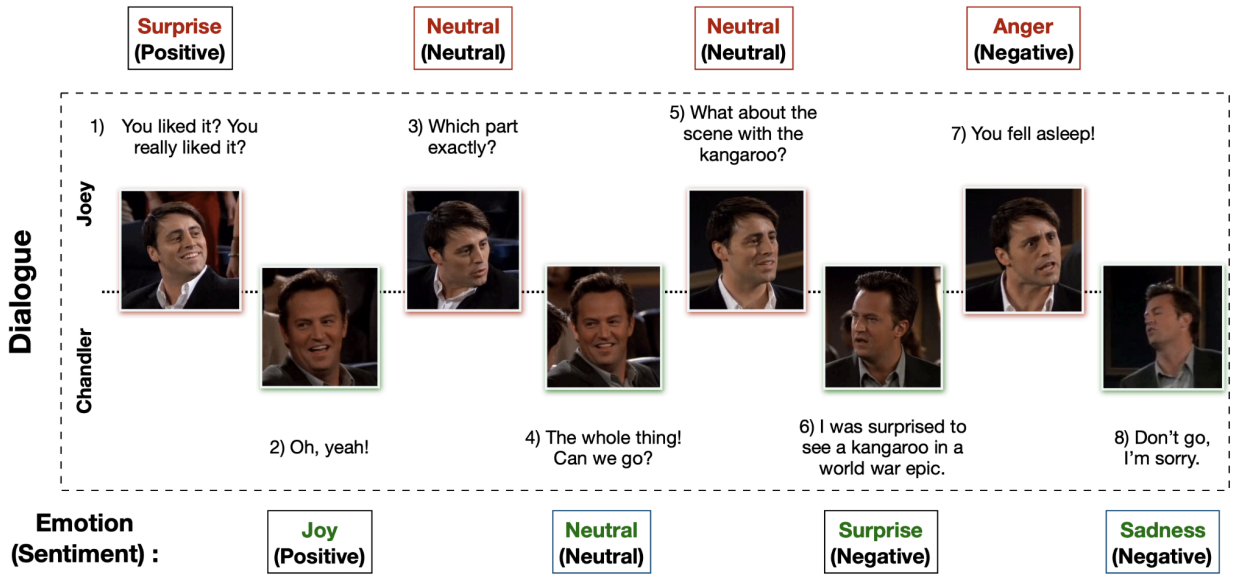

Fig. 3

Emotion recognition in conversation through text involves analyzing written or typed language to identify and understand the emotional states expressed by individuals during a conversation. While traditional emotion recognition often relies on facial expressions, vocal intonations, and other non-verbal cues, text-based emotion recognition focuses on the written words and linguistic patterns used in communication.

Key aspects of emotion recognition in text conversations include:

- Sentiment Analysis: This is a common approach in text-based emotion recognition. Sentiment analysis involves determining the overall sentiment or emotional tone of a

piece of text, whether it be positive, negative, or neutral. Natural language processing (NLP) techniques and machine learning algorithms are often used for sentiment analysis.

- Emotion Classification: Rather than just determining positive or negative sentiment, more advanced systems aim to classify specific emotions expressed in text. This can include emotions such as joy, anger, sadness, surprise, and others. Machine learning models may be trained on labeled datasets to recognize these emotional categories.
- Context Understanding: Emotion recognition in text benefits from an understanding of the context in which the words are used. The same words might convey different emotions depending on the context of the conversation. Advanced models aim to consider context, relationships between words, and linguistic nuances for more accurate emotion recognition.

1.4 Challenges in Emotion Recognition in Conversation

Emotion recognition in conversation, particularly through natural language processing (NLP), faces several challenges due to the complexity and subjectivity of human emotions. Here are some key challenges associated with emotion recognition in conversation:

1. Ambiguity and Nuance: Conversations often involve ambiguous and nuanced expressions of emotions. The same words or phrases can convey different emotions based on context, tone, and individual nuances. Understanding and interpreting these subtleties accurately is challenging for automated systems.
2. Cultural Variability: Emotions are expressed and interpreted differently across cultures. Cultural nuances, idioms, and expressions may not be easily translatable or transferable, making it challenging to develop universal emotion recognition models that perform well across diverse cultural contexts.
3. Context Dependency: The interpretation of emotions heavily depends on the context of the conversation. Words or expressions that seem neutral in isolation may carry emotional connotations within a specific context. Capturing and analyzing contextual information is crucial for accurate emotion recognition.
4. Multi-Modal Challenges: In real-life conversations, emotions are expressed through multiple modalities, including facial expressions, vocal intonation, and body language, in addition to text. Integrating and processing information from these various modalities to provide a holistic understanding of emotions poses additional complexity.
5. Subjectivity and Individual Differences: Emotions are subjective experiences, and individuals may express and perceive them differently. Personal differences, communication styles, and emotional states can vary widely, making it challenging to create models that generalize well across diverse user populations.

6. Temporal Dynamics: Emotions in a conversation can evolve over time. Analyzing the temporal dynamics of emotional expressions and understanding how emotions change during the course of a conversation requires sophisticated modeling techniques.
7. Limited Labeled Datasets: Developing accurate emotion recognition models relies on large and diverse labeled datasets. However, creating such datasets with accurately annotated emotional labels is challenging, and the available datasets may not cover the full spectrum of emotional expressions.
8. Handling Sarcasm and Irony: Conversations often involve the use of sarcasm, irony, or other forms of figurative language. Recognizing these nuances requires a deep understanding of context and cultural references, which can be challenging for automated systems.
9. Privacy and Ethical Considerations: Emotion recognition in conversation raises privacy concerns, as it involves analyzing personal and potentially sensitive information. Ensuring ethical and responsible use of emotion recognition technologies is an ongoing challenge.
10. Real-World Application Challenges: Implementing emotion recognition in real-world applications, such as customer service or mental health support, requires systems that can operate accurately and reliably in dynamic and uncontrolled environments.

Addressing these challenges requires a multidisciplinary approach, combining expertise in NLP, machine learning, psychology, and cultural studies. Ongoing research aims to improve the robustness and accuracy of emotion recognition models in diverse and real-world scenarios.

1.5 Translation

Translation is taking the meaning of a word from the source text and providing an equivalent text in the target language. Many times when you use an online translator, what you are really getting is transliteration or transcription and not a true translation. That is why the results of the "translation" often do not make sense.

It should be noted that there are levels of translation. A simple word-by-word translation takes text from one language and changes it into a word with an identical meaning in the other language. This is great unless there are word strings or sentences in which case a word-for-word translation is not sufficient. For example, the English phrase "under the weather" is used to denote that somebody has fallen ill. The word-to-word Spanish translation for the same would be "por debajoel clima", which does not render any meaning in this specific context and the translated message would make no sense at all in the target language. It did not consider sentence structure, grammar, cultural issues, context, or overall meaning.

To be effective, a translation must take the meaning behind the text and put it into the target language so that the intent of the message remains intact. In the dead meat example, an effective translation might be “sentirse enfermo” or “feeling ill”. A comprehensive translation will consider the grammar, syntax, and local culture during the conversion process so that the final document reads as if a native speaker wrote it.

Translation in the context of Natural Language Processing (NLP) refers to the process of converting text or speech from one language into another. It is a fundamental task in NLP, and its applications are diverse, contributing to a wide range of practical and technological advancements.

1.6 Transliteration

Transliteration is a problem in natural language processing for transforming the text in one script into another script so as to preserve the phonetic structure of words as closely as possible. It is useful when a user knows a language but does not know how to write its script. Machine transliteration can play an important role in natural language processing applications such as information retrieval and machine translation, especially for handling proper nouns and technical terms, cross-language applications, data mining, and information retrieval systems.

Transliteration is the conversion of a word from one language to another without losing its phonological characteristics. It is the practice of transcribing a word or text written in one writing system into another writing system. For instance, the English word school would be transliterated to the Hindi word कूल. Note that this is different from the translation in which the word school would map to पाठशाला (’paathshaala’).

Hence, transliteration can be understood as the process of entering data in one language using the script of another language. In general, the mapping between the alphabet of one language and the other in a transliteration scheme will be as close as possible to the pronunciation of the word. Transliteration is a grapheme-based conversion method that converts every grapheme of the source language to its equivalent in the target language.

2. LITERATURE REVIEW

Recognizing emotions and detecting the triggers for change of emotions in the course of a dialogue narrative or any other means of conversation have lately seen a phase of prime importance in the field of Artificial Intelligence and Natural Language Processing, in particular. This has been primarily due to the high dependence of various research areas on this specialism, ranging from virtual assistants, business survey analysis, market research, law enforcement surveillance, psychological diagnosis and treatment, smart education, etc.

To combat this highly emerging subdomain, several research works have been carried out dealing with the art of detecting various emotions through text and other forms of data, as illustrated in Table 1. For instance, in the research study [1]), an attempt has been made to address the challenges associated with the task of emotion recognition in conversations along with the existing approaches and their drawbacks for doing the same. In [6], a combination of semantic/syntactic and emotional information has been incorporated by adopting pre-trained word representations to recognize emotions from textual data. Another research work [2], involves the task of emotion classification of web blog corpora using support vector machine (SVM) and conditional random field (CRF) machine learning techniques.

In [7], a comparative analysis was performed among BERT, RoBERTa, DistilBERT, and XLNet for text-based emotion recognition. An accuracy of 0.7431, 0.7299, 0.7009, and 0.6693 was achieved for RoBERTa, XLNet, BERT, and DistilBERT (respectively). In the research study [4], a novel approach was proposed with EmoTxt, a toolkit for emotion recognition from text and for the training of custom emotion classification models.

Another research study [5], proposed a multimodal emotion detection framework that extracts multimodal features from conversational videos and hierarchically models the self- and inter-speaker emotional influences into global memories. In a similar study [9], a context- and sentiment-aware framework, termed Sentic GAT, which tends to select the commonsense knowledge consistent with the context semantics and sentiment of the target utterance, was proposed to solve this challenge.

References	**Paper Title**	**Dataset Used**	**Models Used**	**Results Obtained**
[2]	Emotion Classification Using Web Blog Corpora	Blog Posts	SVM and CRF	Refer Table 7 (in publication)
[3]	Multiple Emotions Detection in Conversation Transcripts	Cornell Movie Dialog dataset, ISEAR data	Neural Network	Refer Table 4 (in publication)

[4]	EmoTxt: A Toolkit for Emotion Recognition from Text	Stack Overflow and Jira posts	EmoTxt classification models	Refer Table 2 (in publication)
[5]	ICON: Interactive Conversational Memory Network for Multimodal Emotion Detection	IEMOCAP and SEMAINE	Interactive Conversationa l memory Network (ICON)	Refer Table 6 (in publication)
[6]	Semantic-Emotion Neural Network for Emotion Recognition From Text	Aggregated annotated corpora for Emotion Classification in Text	Semantic-Emo tion Neural Network (biLSTM \& CNN)	Refer Table 3 (in publication)
[7]	Comparative Analyses of BERT, RoBERTa, DistilBERT, and XLNet for Text-based Emotion Recognition	ISEAR data	BERT, RoBERTa, DistilBERT, and XLNet pre-trained transformer models	Refer Table 2 (in publication)
[8]	Contextualized Emotion Recognition in Conversation as Sequence Tagging	IEMOCAP, MELD and Dailydialogue	DialogueRNN	Refer Table 4 (in publication)
[9]	Context- and Sentiment-Aware Networks for Emotion Recognition in Conversation	DailyDialog, MELD and EmoryNLP	Sentic GAT	Refer Table 3 (in publication)

Table 1

In the research work [3], a method of predicting emotions from multi-label conversation transcripts has been developed. The proposed methodology includes building an emotion lexicon bootstrapped from Wordnet following the notion of Plutchik's basic emotions and dyads. A similar study [8], proposes a method to model the ERC task as sequence tagging where a Conditional Random Field (CRF) layer is leveraged to learn the emotional consistency in the conversation. With this, LSTM-based encoders are employed that capture self and inter-speaker dependency of interlocutors to generate contextualized utterance representations which are fed

into the CRF layer. For capturing long-range global context, a multi-layer Transformer encoder has been used to enhance the LSTM-based encoder.

Several research works have been carried out for the task of reasoning of emotion flip (EFR) as well, as mentioned in Table 2. For instance, in [10], the task of Instigator-based Emotion Flip Reasoning (EFR) has been kept forward, which aims to identify the instigator behind a speaker's emotion flip within a conversation. As a part of this work, a novel neural architecture called TGIF has been developed, which leverages Transformer encoders and stacked GRUs to capture the dialogue context, speaker dynamics, and emotion sequence in a conversation.

References	**Paper Title**	**Dataset Used**	**Models Used**	**Results Obtained**
[11]	Emotion Estimation and Reasoning Based on Affective Textual Interaction	WordNet-Affect Database of ITS-irst	Emotion estimation module based on Word-spotting technique	-
[12]	Recognizing Emotion Cause in Conversations	RECCON (built in this work using IEMOCAP and DailyDialog)	RoBERTa Base and Large, SpanBERT Fine-tuned on SQuAD, ECPE-2D, ECPE-MLL, RankCP	Refer Tables 5 and 6 (in publication)
[10]	Emotion Flip Reasoning in Multiparty Conversations	MELD-I	TGIF	Refer Table 4 (in publication)

Table 2

In the research study [11], the methodology of emotion estimation has been proposed that assesses the affective content from textual messages. to detect emotion from chat or other dialogue messages and to employ animated agents capable of emotional reasoning based on textual interaction.

Another research study [12], featured a new annotated dataset for the task of recognizing emotion cause in conversations, RECCON. This study formulated two distinct subtasks of recognizing emotion cause in conversations: (i) causal span extraction and (ii) causal emotion entailment. For the first subtask, RoBERTa Base, and SpanBERT Fine-tuned on SQuAD, were incorporated in this study, while the second subtask involved the use of RoBERTa Base and Large, ECPE-2D, ECPE-MLL, and RankCP.

3. RESEARCH GAPS

- Lexicon-based and modern deep learning-based, vanilla emotion recognition ignores conversation-specific factors such as the presence of contextual cues, the temporality in speakers' turns, or speaker-specific information.
- The state-of-the-art methods keep mimicking the same emotion for a particular party and fail in most cases where a change occurs.
- If two speakers use completely opposite emotions to express their common interest or opinion, it is not possible for a vanilla emotion recognizer to understand the positive emotion of both the interlocutors on the aspect of the opinion.
- Naturally, emotion recognition in multiparty conversations is more challenging in comparison with dyadic conversations due to the difficulty in tracking individual speaker states and handling co-references.
- An ERC system, incapable of detecting sarcasm, mostly fails to predict the emotion of the sarcastic utterances correctly.

4. RESEARCH OBJECTIVES

- To devise a mechanism for interpreting conversation-specific factors, by comprehending contextual cues, and understanding the temporality in speakers' turns, along with the speaker-specific information.
- To implement a system capable of monitoring the emotion sequence, tracking the triggers that can possibly result in a turn of the speaker's emotions.
- To ensure that the proposed system is able to comprehend the assertion and negation aspects resulting from a dialogue.
- To handle multiparty conversations in addition to dyadic conversations, while tracking individual speaker states and handling co-references.
- To prepare a methodology for handling sarcasm-inclusive conversations accurately.

5. THEORETICAL BACKGROUND

5.1 TRANSFORMERS

A transformer model is a neural network that learns context and thus meaning by tracking relationships in sequential data like the words in this sentence. Transformer models apply an evolving set of mathematical techniques, called attention or self-attention, to detect subtle ways even distant data elements in a series influence and depend on each other.

Transformers can translate text and speech in near real-time, opening meetings and classrooms to diverse and hearing-impaired attendees. They're helping researchers understand the chains of genes in DNA and amino acids in proteins in ways that can speed drug design. Transformers can detect trends and anomalies to prevent fraud, streamline manufacturing, make online recommendations, or improve healthcare. People use transformers every time they search on Google or Microsoft Bing.

5.2 TRANSFER LEARNING

Pretraining is the act of training a model from scratch: the weights are randomly initialized, and the training starts without any prior knowledge.

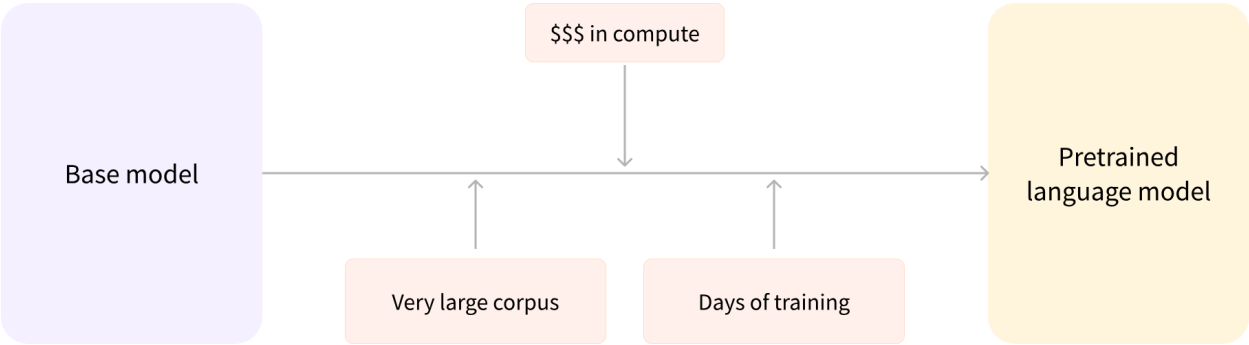

Fig. 4

This pretraining is usually done on very large amounts of data. Therefore, it requires a very large corpus of data, and training can take up to several weeks. Fine-tuning, on the other hand, is the training done after a model has been pre-trained. To perform fine-tuning, you first acquire a pre-trained language model, then perform additional training with a dataset specific to your task.

Fine-tuning a model takes less time, data, financial, and environmental costs. It is also quicker and easier to iterate over different fine-tuning schemes, as the training is less constraining than full pretraining. This process will also achieve better results than training from scratch (unless

you have lots of data), which is why you should always try to leverage a pre-trained model — one as close as possible to the task you have at hand — and fine-tune it.

5.3 ARCHITECTURE OF TRANSFORMERS

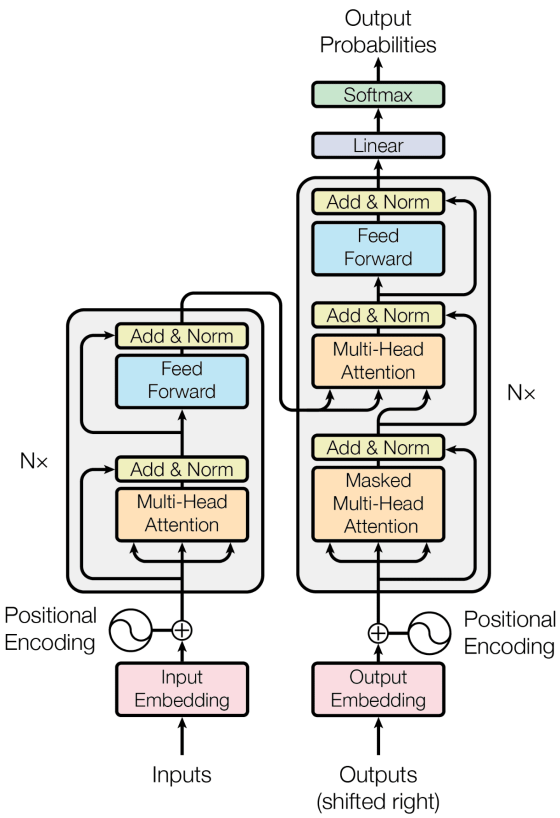

Fig. 5

The Transformer architecture follows an encoder-decoder structure but does not rely on recurrence and convolutions to generate an output. In a nutshell, the task of the encoder, on the left half of the Transformer architecture, is to map an input sequence to a sequence of continuous representations, which is then fed into a decoder. The decoder, on the right half of the architecture, receives the output of the encoder together with the decoder output at the previous time step to generate an output sequence.

Transformers leverage attention mechanisms, multi-head attention, positional encodings, normalization techniques, and layer stacking to efficiently process input sequences, capture dependencies, and generate outputs, allowing them to excel in various natural language processing tasks without relying on recurrent or convolutional structures.

- Attention Mechanism: At the heart of the transformer architecture lies the attention mechanism, which allows the model to weigh different parts of the input sequence differently when processing information. This mechanism enables the model to focus on relevant words or tokens within the sequence during both encoding and decoding stages. Self-attention, a key aspect of the transformer, allows each word in the sequence to attend to all other words, capturing dependencies and relationships effectively without relying on sequential processing.
- Multi-Head Attention: Transformers employ multiple attention heads within each layer, allowing the model to capture different aspects and patterns of relationships within the input sequence. This multi-head attention mechanism allows for the parallel processing of information, enhancing the model's ability to learn intricate patterns and dependencies in the data.

- Positional Encoding: Since transformers lack inherent positional information due to the absence of recurrence or convolutional operations, positional encodings are added to the input embeddings. These encodings provide the model with information about the order of tokens in the sequence, enabling it to understand the sequential order of the input.
- Layer Normalization and Feed-Forward Networks: Each transformer layer consists of sub-layers involving layer normalization and feed-forward neural networks. Layer normalization helps in stabilizing the training process by normalizing the inputs across the feature dimension. Feed-forward networks, with activation functions like ReLU, facilitate non-linear transformations, contributing to the model's ability to learn complex relationships between tokens.
- Residual Connections and Layer Stacking: To facilitate the smooth flow of information and alleviate the vanishing gradient problem, residual connections are employed, allowing gradients to flow more effectively through the network. Additionally, transformers typically stack multiple encoder and decoder layers, each building upon the representations learned by the preceding layers, enabling the model to capture hierarchical and abstract features.

5.4 ATTENTION LAYERS

A key feature of Transformer models is that they are built with special layers called attention layers. Consider the task of translating text from English to French. Given the input “You like this course”, a translation model will need to also attend to the adjacent word “You” to get the proper translation for the word “like”, because in French the verb “like” is conjugated differently depending on the subject. The rest of the sentence, however, is not useful for the translation of that word. In the same vein, when translating “this” the model will also need to pay attention to the word “course”, because “this” translates differently depending on whether the associated noun is masculine or feminine.

Again, the other words in the sentence will not matter for the translation of “this”. With more complex sentences (and more complex grammar rules), the model would need to pay special attention to words that might appear farther away in the sentence to properly translate each word. The same concept applies to any task associated with natural language: a word by itself has a meaning, but that meaning is deeply affected by the context, which can be any other word (or words) before or after the word being studied.

5.5 ENCODER MODELS

Encoder models use only the encoder of a Transformer model. At each stage, the attention layers can access all the words in the initial sentence. These models are often characterized as having “bi-directional” attention and are often called auto-encoding models. The pretraining of these

models usually revolves around somehow corrupting a given sentence (for instance, by masking random words in it) and tasking the model with finding or reconstructing the initial sentence.

Encoder models are best suited for tasks requiring an understanding of the full sentence, such as sentence classification, named entity recognition (and more generally word classification), and extractive question answering.

Representatives of this family of models include:

- ALBERT
- BERT
- DistilBERT
- ELECTRA
- RoBERTa

5.6 DECODER MODELS

Decoder models use only the decoder of a Transformer model. At each stage, for a given word the attention layers can only access the words positioned before it in the sentence. These models are often called auto-regressive models. The pretraining of decoder models usually revolves around predicting the next word in the sentence. These models are best suited for tasks involving text generation.

Representatives of this family of models include:

- CTRL
- GPT
- GPT-2
- Transformer XL

5.7 PRE-TRAINED LANGUAGE MODELS

Pre-trained models (PTMs) for natural language processing (NLP) are deep learning models, such as transformers, that have been trained on large datasets to perform specific NLP tasks. By training on extensive corpora, PTMs can learn universal language representations, which are useful for various downstream NLP tasks such as text summarization, named entity recognition, sentiment analysis, part-of-speech tagging, language translation, sentiment analysis, text generation, information retrieval, text clustering, and many more. This eliminates the need to train a new model from scratch each time. In other words, pre-trained models can be seen as reusable NLP models that developers can use to quickly build NLP applications.

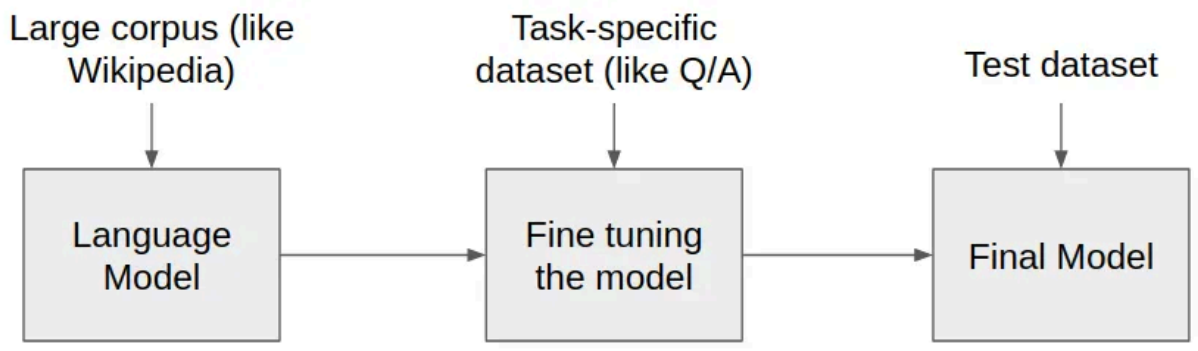

Fig. 6

The language model is first fed a large amount of unannotated data (for example, the complete Wikipedia dump). This lets the model learn the usage of various words and how the language is written in general. The model is now transferred to an NLP task where it is fed another smaller task-specific dataset, which is used to fine-tune and create the final model capable of performing the aforementioned task.

Transformers offer a range of pre-trained deep learning NLP models that cater to different tasks like text classification, question answering, and machine translation. These pre-trained models are freely available and can be utilized without requiring extensive NLP knowledge. While the first generation of PTMs focused on learning good word embeddings, the latest or second generation is designed to learn contextual word embeddings. I'll delve into the details of different types of pre-trained models in my upcoming blogs.

6. PROPOSED METHODOLOGY

6.1 DATASET USED

For the research purpose, we are using a structured and rich collection of conversational data extracted from a Hindi show, organized in the format of JSON data. This dataset is meticulously designed to capture the intricacies of dialogue exchanges, enabling the development of models for emotion analysis. The data is structured into a data frame, a tabular representation commonly used in data science, with four essential columns: episode, speakers, utterances, and emotions.

Dataset Structure:

- Episode Column: In the context of a Hindi show, the 'episode' column serves as a categorical identifier, allowing for the systematic organization of the dataset based on individual episodes. This categorization is crucial for understanding how emotions unfold across different segments of the show, providing a nuanced perspective on the evolution of sentiments throughout the series.
- Speakers Column: The 'speakers' column offers insights into the dynamics of the conversation by documenting the participants involved. This information is vital for discerning the nuances in emotional expression based on the characters engaging in the dialogue. Understanding the context of who is speaking and to whom adds depth to the emotion classification task, enhancing the dataset's potential for nuanced analysis.
- Utterances Column: The 'utterances' column constitutes the core of the dataset, containing the actual text data of the conversations. This text data is the foundation for training models to recognize and classify emotions. The diverse and authentic nature of the utterances ensures that the model is exposed to a wide range of linguistic expressions, making it more robust in its ability to generalize to various conversational styles and contexts.
- Emotions Column: The 'emotions' column is the target variable for machine learning models, assigning specific emotion labels to each utterance. The emotion labels, including disgust, contempt, anger, neutral, joy, sadness, fear, and surprise, cover a comprehensive spectrum of human emotions. This exhaustive set of labels facilitates a nuanced understanding of emotional states within the conversations, allowing the model to distinguish subtle variations in sentiment.

Dataset Size and Split:

The dataset is divided into two sets: the training dataset and the validation dataset. The training dataset comprises 8506 utterances, providing a substantial volume of data for training machine learning models. This dataset's size is conducive to building robust models capable of recognizing patterns and nuances in the expression of emotions in Hindi show conversations. The validation dataset, with 1354 utterances, serves as a separate set to assess the model's

performance on unseen data. This split ensures that the model is tested on different instances than those it encountered during training, promoting generalization and preventing overfitting.

Emotion Labels:
The emotion labels—disgust, contempt, anger, neutral, joy, sadness, fear, and surprise—offer a comprehensive range to capture the diverse spectrum of emotions expressed in the conversations. These labels enable the model to learn and differentiate between nuanced emotional states, enhancing its ability to accurately classify emotions in the dialogue data.

	episode	speakers	utterances	emotions
0	roseshantstobeadopted.wmv	maya	kya kya bhar ke rakha hai indravadan ne is gha...	disgust
1	roseshantstobeadopted.wmv	indu	ok, chalo rosesh chalo bahar	contempt
2	roseshantstobeadopted.wmv	rosesh	momma! hath chhodiye dad!	anger
3	roseshantstobeadopted.wmv	indu	dekho maya ya na bol raha hai! fir tum hi kaho...	neutral
4	roseshantstobeadopted.wmv	maya	indravadan, tum kitne salo se ye joke maar rah...	neutral
...	...	...	...	...
8501	radhabaileaves	rosesh	lekin what about my ghadi? 17000 ki ghadi hai ...	sadness
8502	radhabaileaves	sita or geeta	nhi nhi mujhe yahan pe kaamich nhi karna hai, ...	anger
8503	radhabaileaves	maya	ahan!	joy
8504	radhabaileaves	sita or geeta	goodbye!	fear
8505	radhabaileaves	maya	no goodbye! nikalo! nikalo! chorro hath chorro...	disgust

8506 rows × 4 columns

Fig. 7

6.2 PREPROCESSED DATASET

The preprocessed Data described encompasses seven distinct columns, each serving a specific purpose to enhance the depth and utility of the dataset. The columns are episode, speakers, utterances, emotions, cleaned_utterances, hindi_utterances, and english_utterances.

1. Episode and Speakers Columns
2. Utterances and Emotions Columns
3. Cleaned Utterances Column
4. Hindi Utterances Column
5. English Utterances Column

The inclusion of both Hindi and English versions of the utterances broadens the dataset's applicability. The described preprocessing steps highlight a thoughtful approach to preparing the dataset for analysis and modeling. By incorporating transliteration and translation techniques, the data accommodates a broader audience and facilitates a more comprehensive exploration of emotions and linguistic nuances within the context of Hindi show dialogues. This enriched dataset is poised to contribute to the advancement of natural language processing applications, especially those focusing on multilingual and cross-cultural dimensions.

	episode	speakers	utterances	emotions	cleaned_utterances	hindi_utterances	english_utterances
0	roseshantstobeadopted.wmv	maya	kya kya bhar ke rakha hai indravadan ne is gha...	disgust	kya kya bhar ke rakha hai indravadan ne ghar m...	क्य क्य भर् के रख है इन्द्रवदन् ने इस् घर् मेइ...	What does Indravadan keep in the house? Indrav...
1	roseshantstobeadopted.wmv	indu	ok, chalo rosesh chalo bahar	contempt	ok chalo rosesh chalo bahar	ओक्, चलो रोसेश् चलो बहर्	ok, let's rise, let's spring
2	roseshantstobeadopted.wmv	rosesh	momma! hath chhodiye dad!	anger	momma hath chhodiye dad	मोम्मा हथ् छोदिये दद्!	Momma! Hath Chodiye Dad!
3	roseshantstobeadopted.wmv	indu	dekho maya ya na bol raha hai! fir tum hi kaho...	neutral	dekho maya ya na bol raha hai fir tum hi kahog...	देखो मय य न बोल् रह है! फिर् तुम् हि कहोगि फल्...	Look Maya or I am not speaking! Then you will ...
4	roseshantstobeadopted.wmv	maya	indravadan, tum kitne salo se ye joke maar rah...	neutral	indravadan tum kitne salo se ye joke maar rahe...	इन्द्रवदन्, तुम् किल्ने सलो से ये जोके मार् रह...	Indravadan, for how many years have you been t...
...	...	...	...	...	...	...	...
8501	radhabaileaves	rosesh	lekin what about my ghadi? 17000 ki ghadi hai ...	sadness	lekin ghadi ki ghadi hai momma recurring alarm...	लेकिन् व्हत् अबोउत् म्य् घदि? १७००० कि घदि है ...	but what about my Ghadi? 17000 ki ghadi hai mo...
8502	radhabaileaves	sita or geeta	nhi nhi mujhe yahan pe kaamich nhi karna hai, ...	anger	nhi nhi mujhe yahan pe kaamich nhi karna hai j...	न्हि न्हि मुझे यहन् पे कामिच् न्हि वर्न है, ने...	No no, I don't want to do evil here, I am here!
8503	radhabaileaves	maya	ahan!	joy	ahan	अहन्!	oh!
8504	radhabaileaves	sita or geeta	goodbye!	fear	goodbye	गूद्ब्ये!	goodbye!
8505	radhabaileaves	maya	no goodbye! nikalo! nikalo! chorro hath chorro...	disgust	goodbye nikalo nikalo chorro hath chorro badby...	नो गूद्ब्ये! निकलो! निकलो! चोर्रो हथ् चोर्रो, ...	no goodbye! take it out! take it out! chorro h...

8506 rows × 7 columns

Fig. 8

6.3 IMPLEMENTATION STRATEGY

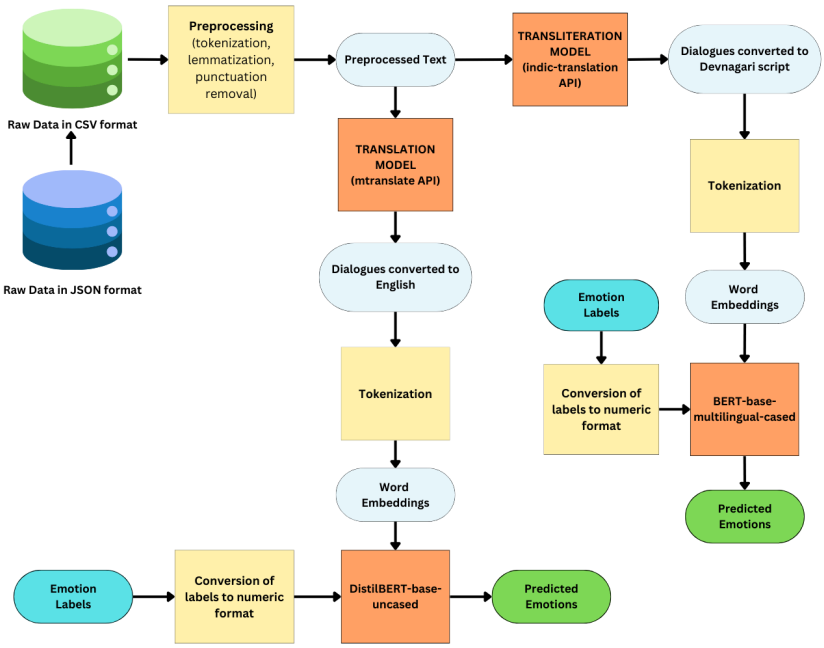

Fig. 9

1. Preprocessing the JSON file:
    a. Read the JSON file
    b. Tokenization: Use libraries like NLTK (Natural Language Toolkit) or SpaCy to split the sentences into words (tokenization).
    c. Lemmatization: Apply lemmatization to the words obtained from the tokenization step. NLTK or SpaCy can also help perform lemmatization.
    d. Punctuation Removal: Utilize string manipulation methods or regular expressions to remove punctuation from the text.
2. Transliteration of code-mixed dialogues to Hindi script:

    a. Indic-Translation API: Make API requests using appropriate libraries (e.g., requests in Python) to the indic-translation API. Pass the code-mixed dialogues as input and retrieve the transliterated text in Hindi.
3. Translation of code-mixed dialogues to English:
    a. mtranslate API: Use the mtranslate API through API requests to translate code-mixed dialogues to English. Input the dialogues and receive the translated text in English.
4. Conversion of labels to numeric format:
    a. Mapping Labels: Create a mapping dictionary or function to convert the existing labels to numeric values (0-7). For instance, if the labels are strings, assign each unique label a numeric value from 0 to 7.

Fine-tuning models:

5. For English Utterances (DistilBERT-base-uncased):
    a. DistilBERT Model: Utilize the Hugging Face Transformers library to access the DistilBERT model for English text.
    b. Tokenization: Tokenize the English utterances using the DistilBERT tokenizer to obtain encodings suitable for the model.
    c. Fine-tuning: Fine-tune using the specified training arguments (batch size, learning rate, epochs, etc.) by training the DistilBERT model on your dataset.
    d. Softmax Activation: Apply the softmax activation function to the logits (output scores) obtained from the model after inference.
    e. Store Predictions: Save the predictions along with relevant information like dialogue IDs, timestamps, or any other necessary details.
6. For Hindi Utterances (BERT-base-multilingual-cased):
    a. BERT Model: Use the Hugging Face Transformers library to access the BERT model for Hindi text.
    b. Tokenization: Tokenize the Hindi utterances using the BERT tokenizer to obtain encodings suitable for the model.
    c. Fine-tuning: Fine-tune the BERT model on your dataset using the specified training arguments (batch size, learning rate, epochs, etc.).
    d. Softmax Activation: Apply softmax activation to the logits obtained from the BERT model after inference.
    e. Store Predictions: Save the predictions with necessary details similar to the English model.

6.4 MODEL ARCHITECTURE

BERT multilingual base (cased)

BERT, an unsupervised model, utilizes the Masked Language Model (MLM) and Next Sentence Prediction (NSP) tasks during training. The MLM task involves masking random words in sentences and predicting them based on context, while NSP trains the model to understand relationships between sentence pairs. This dual-task approach fosters a deeper understanding of contextual nuances. The model architecture consists of multiple Transformer layers, enabling efficient computation of contextual embeddings. BERT's attention mechanism allows it to capture intricate relationships among words in both directions, thereby encoding rich contextual information. Furthermore, BERT's tokenization method uses WordPiece embeddings, breaking down words into subword units to handle out-of-vocabulary words and improve representation learning. The multilingual base model's training corpus encompasses a wide range of languages, enabling it to learn language-agnostic features and facilitating transfer learning across diverse language sets. This capability makes it highly adaptable for tasks requiring multilingual understanding, such as machine translation, sentiment analysis, and cross-lingual information retrieval.

DistilBERT base (uncased)

DistilBERT base (uncased) is a compressed version of BERT, employing knowledge distillation to condense BERT's essence into a smaller architecture, balancing size and performance. This model aims to widen accessibility for applications facing resource limitations in natural language processing (NLP). Utilizing knowledge distillation, DistilBERT is trained to replicate BERT's output probabilities while employing fewer parameters, resulting in a more compact yet expressive model. The reduced size offers advantages such as decreased computational resources, facilitating deployment in resource-constrained scenarios and real-time applications. Its efficiency in inference speed makes DistilBERT suitable for production environments and devices with limited processing capabilities. Despite its smaller footprint, it maintains competitive performance across diverse NLP benchmarks, striking a balance between computational efficiency and task-specific effectiveness. DistilBERT is favored by practitioners seeking efficient yet high-performing models for NLP applications, particularly beneficial when computational resources are restricted while aiming to achieve the advantages associated with larger models like BERT.

6.5 MODEL TRAINING

DistilBERT base (uncased)

- Custom Model Definition: DistilBertForMultilabelSequenceClassification
    - This class is derived from DistilBertForSequenceClassification, which is a sequence classification model based on DistilBERT.
- Initialization (__init__):
    - Initializes the custom model by calling the superclass constructor (super().__init__(config)).
- Forward Method (forward):
    - Accepts various input arguments (input_ids, attention_mask, etc.).
    - Passes inputs through the DistilBERT model (self.distilbert) and retrieves the hidden states.
    - Extracts the pooled output (representation of the entire input sequence) from the hidden states.
    - Applies dropout to the pooled output for regularization.
    - Feeds the output through a classifier layer (self.classifier) to obtain logits.
    - Calculates the loss using Binary Cross Entropy (BCEWithLogitsLoss) if labels are provided.
    - Returns a SequenceClassifierOutput object containing loss, logits, hidden states, and attentions if return_dict is set or a tuple otherwise.
- Training Configuration: TrainingArguments
    - output_dir: Specifies the directory to save the trained model and associated files.
    - evaluation_strategy: Defines when to perform evaluation during training ("epoch" here means at the end of each epoch).
    - save_strategy: Indicates how model checkpoints should be saved ("epoch" here means one checkpoint per epoch).
    - learning_rate: Sets the learning rate for training the model.
    - per_device_train_batch_size: Specifies the batch size per GPU during training.
    - per_device_eval_batch_size: Specifies the batch size per GPU during evaluation.
    - num_train_epochs: Specifies the number of epochs for training.
    - weight_decay: Defines the weight decay rate for controlling overfitting.
    - logging_steps: Determines the frequency of logging training loss.
- Training Process using Trainer
    - Trainer from the transformers library is used to facilitate the training loop.
    - It takes in the defined model, training arguments (args), and other necessary inputs (like datasets) to execute the training process.
    - The Trainer object orchestrates the training loop, including forward passes, loss computation, backward propagation (if labels are available), and parameter updates based on the optimizer and learning rate schedule.

BERT multilingual base (cased)

- Training Configuration:
  The given TrainingArguments define the settings and hyperparameters for training the BERT-multilingual model.
  - output_dir: Specifies the directory to save the trained model, predictions, and checkpoints.
  - num_train_epochs: Sets the number of training epochs to 3, controlling the number of times the entire dataset will be iterated over for training.
  - per_device_train_batch_size: Determines the batch size per device (in this case, 32) during training.
  - per_device_eval_batch_size: Specifies the batch size per device (also 32) during evaluation.
  - weight_decay: Defines the weight decay rate (set to 0.01) used as a regularization technique to prevent overfitting.
  - evaluation_strategy: Indicates when to perform evaluation during training ("epoch" means at the end of each epoch).
  - save_strategy: Defines the strategy for saving model checkpoints ("epoch" here means saving one checkpoint per epoch).
  - logging_steps: Determines the frequency of logging training loss. It's calculated as len(train_dataloader) // 32.
- Training Process:
  The provided configuration is used to train a BERT-multilingual model using the Hugging Face Trainer module. The training process involves the following steps:
  - Initialization: Initialize the BERT-multilingual model for training.
  - Data Loading: Load the training dataset (train_dataloader) that contains the necessary data for training the model.
  - Training Loop: Utilize the Trainer object to manage the training loop.
  - Iterate through the dataset for the specified number of epochs (3 epochs in this case).
  - During each epoch:
  - Perform forward passes through the model.
  - Calculate the loss based on the predicted outputs and ground truth labels (if available).
  - Backpropagate the gradients and update the model's parameters.
  - Log the training loss at specified intervals (logging_steps).
  - Evaluate the model's performance on the evaluation dataset at the end of each epoch based on the specified evaluation strategy.
  - Save model checkpoints at intervals determined by save_strategy.

7. EVALUATION METRICS

7.1 DistilBERT BASE (UNCASED)

[798/798 04:01, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy Thresh
1	0.315500	0.288009	0.887094
2	0.279600	0.281223	0.888109
3	0.263100	0.279579	0.887925

[43/43 01:03]

```
TrainOutput(global_step=798, training_loss=0.2858860726493941, metrics={'train_runtime': 241.4669,
'train_samples_per_second': 105.679, 'train_steps_per_second': 3.305, 'total_flos':
394383152244096.0, 'train_loss': 0.2858860726493941, 'epoch': 3.0})
```

Fig. 10

	precision	recall	f1-score	support
anger	0.4250	0.1441	0.2152	118
contempt	0.1250	0.0270	0.0444	74
disgust	1.0000	0.0476	0.0909	21
fear	0.0000	0.0000	0.0000	88
joy	0.4673	0.4079	0.4356	228
neutral	0.5188	0.8515	0.6447	633
sadness	0.2632	0.0397	0.0690	126
surprise	0.3750	0.2273	0.2830	66
accuracy			0.4963	1354
macro avg	0.3968	0.2181	0.2229	1354
weighted avg	0.4234	0.4963	0.4176	1354

Fig. 11

7.2 BERT MULTILINGUAL BASE (CASED)

[1905/1905 54:07, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy	F1	Precision	Recall
1	1.261800	2.206405	0.459558	0.078715	0.057445	0.125000
2	1.241200	2.365365	0.459558	0.078715	0.057445	0.125000
3	1.218600	2.343876	0.459558	0.078715	0.057445	0.125000

Fig. 12

	precision	recall	f1-score	support
anger	0.0000	0.0000	0.0000	118
contempt	0.0000	0.0000	0.0000	74
disgust	0.0000	0.0000	0.0000	21
fear	0.0000	0.0000	0.0000	88
joy	0.0000	0.0000	0.0000	228
neutral	0.4675	1.0000	0.6371	633
sadness	0.0000	0.0000	0.0000	126
surprise	0.0000	0.0000	0.0000	66
accuracy			0.4675	1354
macro avg	0.0584	0.1250	0.0796	1354
weighted avg	0.2186	0.4675	0.2979	1354

Fig. 13

8. ANALYSIS OF RESULTS

8.1 DistilBERT BASE (UNCASED)

The provided training and validation results pertain to the training of a DistilBERT-base-uncased model for emotion recognition in multiparty conversation. Let's delve into the details of these results to gain a deeper understanding of the model's performance over the three epochs:

- Epoch 1:
    - Training Loss (Epoch 1): 0.315500
    - Validation Loss (Epoch 1): 0.288009
    - Accuracy Threshold (Epoch 1): 0.887094

  In the first epoch, the model achieved a training loss of 0.315500, indicating how well it fit the training data. Simultaneously, the validation loss was 0.288009, representing the model's performance on unseen validation data. The accuracy threshold at 0.887094 indicates the level of accuracy achieved by the model on the validation set, where predictions surpassing this threshold are considered correct.
- Epoch 2:
    - Training Loss (Epoch 2): 0.279600
    - Validation Loss (Epoch 2): 0.281223
    - Accuracy Threshold (Epoch 2): 0.888109

  In the second epoch, there is a reduction in both training and validation losses. The training loss decreased to 0.279600, suggesting that the model is improving its ability to fit the training data. The validation loss also decreased to 0.281223, indicating improved generalization to unseen data. The accuracy threshold slightly increased to 0.888109, suggesting that the model's predictions on the validation set are now considered accurate at a slightly higher threshold.
- Epoch 3:
    - Training Loss (Epoch 3): 0.263100
    - Validation Loss (Epoch 3): 0.279579
    - Accuracy Threshold (Epoch 3): 0.887925

  By the third epoch, both training and validation losses continue to decrease. The training loss is now 0.263100, indicating further refinement of the model's representation of the training data. The validation loss is 0.279579, suggesting improved generalization to unseen data. The accuracy threshold at 0.887925 indicates the model's accuracy on the validation set.

These results collectively demonstrate the model's learning progression over the three epochs. The decreasing training and validation losses suggest that the model is effectively capturing patterns in the training data and generalizing well to unseen data. The accuracy threshold

provides a threshold for considering model predictions as accurate, and its variations across epochs highlight the trade-off between precision and recall.

Parameters:

```
TrainOutput(
  global step-798,
  training loss-0.2858860726493941,
  metrics=[
    'train runtime: 241.4669',
    'train samples per second: 105.679',
    'train steps per second: 3.305',
    'total flos: 394383152244096.0',
    'train_loss: 0.2858860726493941',
    'epoch: 3.0'
  ]
)
```

8.2 BERT MULTILINGUAL BASE (CASED)

- Epoch 1:
  - Training Loss (Epoch 1): 1.261800
  - Validation Loss (Epoch 1): 2.206405
  - Accuracy (Epoch 1): 0.459558
  - F1 Score (Epoch 1): 0.078715
  - Precision (Epoch 1): 0.057445
  - Recall (Epoch 1): 0.125000

  In the first epoch, the training loss is 1.261800, indicating how well the model is fitting the training data. The validation loss is 2.206405, representing the model's performance on unseen validation data. The accuracy achieved on the validation set is 0.459558. The F1 score, precision, and recall metrics provide additional insights into the model's performance, considering both false positives and false negatives.
- Epoch 2:
  - Training Loss (Epoch 2): 1.241200
  - Validation Loss (Epoch 2): 2.365365
  - Accuracy (Epoch 2): 0.459558
  - F1 Score (Epoch 2): 0.078715
  - Precision (Epoch 2): 0.057445
  - Recall (Epoch 2): 0.125000

  In the second epoch, there is a slight reduction in the training loss to 1.241200, indicating further improvement in the model's fit to the training data. However, the validation loss

increases to 2.365365, suggesting that the model's generalization to unseen data may have slightly degraded. The accuracy and other metrics remain constant, indicating consistent model performance.

- Epoch 3:
    - Training Loss (Epoch 3): 1.218600
    - Validation Loss (Epoch 3): 2.343876
    - Accuracy (Epoch 3): 0.459558
    - F1 Score (Epoch 3): 0.078715
    - Precision (Epoch 3): 0.057445
    - Recall (Epoch 3): 0.125000

  In the third epoch, there is a further reduction in the training loss to 1.218600, suggesting continued improvement in the model's fit to the training data. The validation loss remains high at 2.343876, indicating challenges in generalizing to unseen data. The accuracy and other metrics stay consistent across epochs.

Overall, the results suggest that, while the training loss decreases, the model's generalization to the validation set remains relatively constant, as evidenced by the consistent accuracy, F1 score, precision, and recall values across the three epochs. It's important to carefully monitor validation metrics to ensure that the model is not overfitting to the training data and is capable of making accurate predictions on new, unseen data.

9. FUTURE WORK

1. Advanced Algorithm Development:
    a. Model Enhancement: Explore and integrate more sophisticated machine learning and NLP techniques like transformer-based architectures (BERT, GPT) or domain-specific models to improve trigger identification accuracy.
    b. Multi-modal Analysis: Incorporate audio-visual analysis alongside textual data to capture a comprehensive understanding of emotional triggers.
2. Emotion Recognition and Contextual Analysis:
    a. Nuanced Emotion Detection: Integrate models capable of recognizing nuanced emotions beyond basic sentiments, enabling the system to understand complex emotional states.
    b. Contextual Understanding: Develop mechanisms to analyze historical context, cultural nuances, and individual characteristics to better interpret triggers within a given context.
3. Data Enrichment and Annotation:
    a. Human Annotated Data: Collect additional annotated data from human annotators to expand and refine the training dataset, focusing on labeling triggers and emotional changes in dialogues.
    b. User Feedback Integration: Incorporate mechanisms for users to provide feedback on trigger identification to continuously enhance model accuracy.
4. Real-time Monitoring and Intervention:
    a. Real-time Analysis: Develop real-time monitoring capabilities to detect triggers and emotional changes during ongoing conversations.
    b. Intervention Strategies: Implement systems that not only identify triggers but also suggest or implement intervention strategies to manage emotional situations.
5. Ethical and Privacy Considerations:
    a. Ethical Compliance: Ensure compliance with ethical guidelines and regulations related to user privacy and consent in emotionally sensitive data analysis.
    b. Privacy Protection: Strengthen data protection measures to safeguard user privacy and transparency in data usage.
6. Cross-lingual and Cross-cultural Adaptability:
    a. Multilingual Analysis: Extend the system's capabilities to detect triggers and emotional changes across different languages.
    b. Cultural Sensitivity: Account for cultural variations in emotional expressions to create a more inclusive and culturally aware model.

10. CONCLUSION

Emotion Discovery and Reasoning its Flip in Conversation (EDiReF) project has addressed the critical challenge of accurately recognizing emotions in sequential dialogues, acknowledging their pivotal role in shaping human communication. Through the implementation of a meticulous mechanism leveraging pre-trained language models such as Multilingual BERT and DistilBERT, this research has made significant strides in understanding speaker-oriented information, dynamics, and crucial features influencing emotional situations within conversations.

The focus on Emotion Recognition in Conversations (ERC) and Emotion Flip Reasoning (EFR) has led to the development of a robust system capable of discerning the nuanced variations in emotions within dialogues. The system's architecture not only interprets the context and sentiments embedded in conversations but also endeavors to detect triggers that catalyze significant emotional shifts among involved parties. The utilization of a code-mixed dataset sourced from a TV show has provided a diverse and comprehensive foundation for this research.

Moving forward, the findings and methodology of this project hold promising implications across various domains, including but not limited to natural language processing, psychology, and human-computer interaction. The proposed mechanism serves as a fundamental stepping stone towards enhancing communication systems by fostering a deeper understanding of emotional cues and triggers, ultimately contributing to more empathetic and contextually aware conversational agents. As the exploration of emotion recognition and reasoning in conversations continues to evolve, this research provides a robust framework for further advancements in this burgeoning field.

11. PROGRAMMING ENVIRONMENT AND TOOLS USED

- System Environment: Python, Google Colab/Jupyter Notebook, GitHub
- Frameworks/Libraries: PyTorch, NLTK, Transformers, etc.
- Models:
    - Transliteration Model(indic-translation API),
    - Translation Model(mtranslate API),
    - Large Language Models:
        - DistilBERT-base-uncased, BERT-base-multilingual-cased.

12. REFERENCES

[1] Soujanya Poria et al. “Emotion recognition in conversation: Research challenges, datasets, and recent advances”. In: IEEE Access 7 (2019), pp. 100943–100953.

[2] Changhua Yang, Kevin Hsin-Yih Lin, and Hsin-Hsi Chen. “Emotion classification using web blog corpora”. In: IEEE/WIC/ACM International Conference on Web Intelligence (WI’07). IEEE. 2007, pp. 275–278.

[3] Duc-Anh Phan, Hiroyuki Shindo, and Yuji Matsumoto. “Multiple emotions detection in conversation transcripts”. In: Proceedings of the 30th Pacific Asia Conference on Language, Information and Computation: Oral Papers. 2016, pp. 85–94.

[4] Fabio Calefato, Filippo Lanubile, and Nicole Novielli. “Emotxt: a toolkit for emotion recognition from text”. In: 2017 seventh international conference on Affective Computing and Intelligent Interaction Workshops and Demos (ACIIW). IEEE. 2017, pp. 79–80.

[5] Devamanyu Hazarika et al. “Icon: Interactive conversational memory network for multimodal emotion detection”. In: Proceedings of the 2018 conference on empirical methods in natural language processing. 2018, pp. 2594–2604.

[6] Erdenebileg Batbaatar, Meijing Li, and Keun Ho Ryu. “Semantic-emotion neural network for emotion recognition from text”. In: IEEE access 7 (2019), pp. 111866–111878.

[7] Acheampong Francisca Adoma, Nunoo-Mensah Henry, and Wenyu Chen. “Comparative analyses of bert, roberta, distilbert, and xlnet for text-based emotion recognition”. In: 2020 17th International Computer Conference on Wavelet Active Media Technology and Information Processing (ICCWAMTIP). IEEE. 2020, pp. 117–121.

[8] Yan Wang et al. “Contextualized emotion recognition in conversation as sequence tagging”. In: Proceedings of the 21th annual meeting of the special interest group on discourse and dialogue. 2020, pp. 186–195.

[9] Geng Tu et al. “Context-and sentiment-aware networks for emotion recognition in conversation”. In: IEEE Transactions on Artificial Intelligence 3.5 (2022), pp. 699–708.

[10] Shivani Kumar et al. “Emotion Flip Reasoning in Multiparty Conversations”. In: arXiv preprint arXiv:2306.13959 (2023).

[11] Chunling Ma, Helmut Prendinger, and Mitsuru Ishizuka. “Emotion estimation and reasoning based on affective textual interaction”. In: Affective Computing and Intelligent Interaction: First International Conference, ACII 2005, Beijing, China, October 22-24, 2005. Proceedings 1. Springer. 2005, pp. 622–628.

[12] Soujanya Poria et al. “Recognizing emotion cause in conversations”. In: Cognitive Computation 13 (2021), pp. 1317–1332.